

Medical Statistics – Lab 7

R version

Part 1: Risk of in-hospital death in patients with acute myocardial infarction

In part 1 of the lab, we are going to analyze the risk of in-hospital death in patients hospitalized because of acute myocardial infarction. The dataset comes from the Worcester Heart Attack Study (WHAS) and includes data from 500 patients admitted in Worcester, Massachusetts in 1997, 1999, and 2001 (file `whas500.sav` from the Datasets menu). The outcome of interest is in-hospital death, measured by the variable “discharge status from hospital” (`dstat`) with values `alive` and `death`.

```
library(haven)    # for reading SPSS files
library(dplyr)    # for data manipulation
library(DescTools) # for performing Hosmer-Lemeshow Goodness of Fit Tests

# Load the dataset
whas500 <- read_sav("datasets/whas500.sav")

# Convert labeled variables to factors
whas500 <- whas500 %>%
  mutate(across(where(is.labelled), as_factor))
```

Exploratory data analysis

To explore whether gender has an effect on the risk of in-hospital death, we start by creating a contingency table and use the table to calculate the proportion in-hospital death in the two gender subgroups:

```
contingencyTable <- table(whas500$gender, whas500$dstat) # create 2 x 2 table
contingencyTable

prop.table(contingencyTable, margin = 1) # Calculate the row proportions (margin = 1)
```

Question 1

Based on the group proportions, do you expect gender to have an effect on the risk of in-hospital death?

Simple logistic regression

To determine whether gender is significantly associated with in-hospital death, we can conduct several statistical tests. As a recap of lab 4, we start by performing a chi-square test of homogeneity.

Question 2

Perform the chi-square test of homogeneity (see instructions in lab 4 if needed). What conclusion can be drawn from the test?

Another option is to perform logistic regression. In R, this can be achieved using the `glm()` function:

```
# Create a 0/1 numeric version of the dependent variable, where a value of 1 means "success"
whas500$dstat_numeric <- ifelse(whas500$dstat=="dead", 1, 0)

# Fit logistic regression model
model.sex <- glm(dstat_numeric ~ gender, family = binomial, data = whas500)
summary(model.sex)
```

Explanation: `glm` stands for “generalized linear model,” an extension of linear regression that accommodates different types of outcome distributions and includes a link function to model the relationship between predictors and the outcome. The argument `family = binomial` specifies that the outcome distribution is Bernoulli/binomial. By default, the link function for this model is “logit,” which corresponds to logistic regression.

Question 3

What is the odds ratio for in-hospital death for females compared to males? How should this odds ratio be interpreted in the context of the study?

Question 4

Based on the estimated regression coefficients (ignoring p-values), what are the predicted proportions of in-hospital deaths for male and female patients? Compare the predicted proportions to the observed proportions from the previously constructed contingency

table. Do they match?

Question 5

What conclusion can be drawn from the logistic regression analysis regarding the association between gender and in-hospital death? Is this in line with the conclusion drawn from the chi-square test?

Multiple logistic regression

To assess the extent to which the effect of gender is confounded by age, we will fit a multiple regression model with in-hospital death as the dependent variable and gender and age as the independent variables:

```
# Fit logistic regression model
model.sex.age <- glm(dstat_numeric ~ gender + age, family = binomial, data = whas500)
summary(model.sex.age)
```

Question 6

How does adjusting for age affect the estimated odds ratio for in-hospital death for females compared to males?

Question 7

Calculate the odds ratio for in-hospital death corresponding to a 10-year increase in age and interpret its meaning.

Likelihood ratio tests

As explained in the syllabus, the p-values in the table of estimated regression coefficients are derived from Wald tests, which test the null hypothesis that the corresponding regression coefficient is equal to 0

Instead of the Wald tests, we can also obtain p-values using likelihood ratio tests, which compare the goodness of fit of the full model (including the predictor of interest) to a reduced model (excluding the predictor) to test the null hypothesis that the predictor has no effect on the outcome. This approach is particularly useful for testing predictors with non-linear or complex effects, as it evaluates their contribution to the model as a whole. Examples include categorical variables with three or more categories (requiring the creation of multiple dummy

variables) and relationships modeled using multiple terms, such as including both a linear and a quadratic term to capture a quadratic relationship.

For example, we can use a likelihood ratio test (LRT) to compare a full model including both sex and age as predictors to a reduced model including only age. To achieve this, we first need to fit the reduced model that only includes age and then perform the LRT using the `anova()` function:

```
# Fit reduced model
model.age <- glm(dstat_numeric ~ age, family = binomial, data = whas500)

# Perform likelihood ratio test
anova(model.age, model.sex.age, test="LRT")
```

Question 8

How does the p-value from the likelihood ratio test compare to the one from the Wald test?

Evaluating model fit

One way to examine the fit of the logistic regression model is the Hosmer-Lemeshow Goodness of Fit test. In R, this test can be performed using the `HosmerLemeshowTest()` function from the `DescTools` package:

```
HosmerLemeshowTest(fit=fitted(model.sex.age), obs=model.sex.age$y, ngr = 10, verbose = TRUE)
```

Explanation:

- `fitted(model.sex.age)` returns the calculated in-hospital death probabilities based on the fitted logistic regression model for each patient in the dataset. These are the
- `model.sex.age$y` retrieves the outcome variable from the fitted logistic regression model.
- `ngr = 10` specifies the number of groups to be used in the test, which is set to 10 by default.
- `verbose = TRUE` not only prints the test-statistic, degrees of freedom, and p-value but also the groups used to perform the test.

The Hosmer-Lemeshow goodness-of-fit test has two variations: the C statistic and the H statistic, which differ in how they group predicted probabilities for comparison. The C statistic groups predicted probabilities into deciles (10 equal-sized groups based on the range of probabilities), while the H statistic uses fixed cutoffs (e.g., evenly spaced intervals between 0 and 1). The C statistic is the most commonly used version as it adapts to the data distribution,

ensuring well-populated groups, making it suitable for general model fit evaluation. In this lab, we therefore focus on this latter statistics (with matches the explanation in the syllabus).

Question 9

Based on the results of the Hosmer-Lemeshow goodness-of-fit test, does our model provide a satisfactory fit to the data?

Part 2: unguided exercises

Exercise 1

Multiple logistic regression was used to construct a prognostic index to predict coronary artery disease from data on 348 patients with valvular heart disease who had undergone routine coronary arteriography before valve replacement (Ramsdale et al. 1982). The estimated equation was:

$$\text{logit}(p) = \ln(p/(1-p)) = b_0 + 1.167 \times x_1 + 0.0106 \times x_2 + \text{other terms}$$

where x_1 stands for the family history of ischaemic disease (0=no, 1=yes) and x_2 is the estimated total number of cigarettes ever smoked in terms of thousand cigarettes, calculated as the average number smoked annually times the number of years smoking.

- What is the estimated odds ratio for having coronary artery disease for subjects with a positive family history relative to subjects with a negative family history?
- What total number of cigarettes ever smoked carries the same risk as a positive family history? Convert the result into years of smoking 20 cigarettes per day.
- What is the odds ratio for coronary artery disease for someone with a positive family history who had smoked 20 cigarettes a day for 30 years compared to a non smoker with no family history?

Exercise 2

Data from 37 patients receiving a non-depleted allogenic bone marrow transplant were examined to see which variables were associated with the occurrence of acute graft-versus-host disease (GvHD: 0=no, 1=yes) (Bagot et al., 1988). Possible predictors are TYPE (type of leukemia: 1=AML, acute myeloid leukaemia; 2=ALL, acute lymphocytic leukaemia; 3=CML, chronic myeloid leukemia), PREG (donor pregnancy: 0= no, 1=yes), and LOGIND (the logarithm of an index of mixed epidermal cell-lymphocyte reactions). The data are in the file `GvHD.sav` available from the Downloads menu.

- (a) Perform a likelihood ratio test to determine whether there is a significant association between the type of leukemia and the occurrence of GvHD after adjusting for donor pregnancy and the logarithm of an index of mixed epidermal cell-lymphocyte reactions.
- (b) In the adjusted model, What is the estimated odds ratio for the occurrence of GvHD for patients with ALL compared to those with ALM?
- (c) Use the Hosmer-Lemeshow goodness-of-fit test to evaluate the fit of the model. Based on the results, does the model provide a satisfactory fit to the data?